

Playing Jedi Academy with Deep Curiosity

Dmitri Tkatch

Abstract

A curious artificial intelligence learns to explore free-for-all maps and defend itself with a saber against built-in bots in duel mode in a few hours of learning on a modest 8GB GPU. It is based on six deep learning neural networks – a future predictor, an action predictor, a state embedder, an image embedder, a snapshot embedder and an A3C actor. The AI can use either raw game data, raw pixels or both as input. It uses a novel form of curiosity where unpredictable futures and predictable actions are rewarded. It defeats Cultist 3-1, Jedi Trainer 1-0 while starting at a disadvantage and the final boss of the game Desann 20-11. We note the importance of curiosity for maintaining a variety of play and our novel form of curiosity for maintaining smooth actions.

Videos – <https://www.youtube.com/@eternalyze0>.

Program – <https://github.com/Eternalyze0/jediknight>.

Correspondence – eternalyze0@gmail.com

1 Introduction

Star Wars Jedi Knight Jedi Academy is a 3D competitive saber-combat game where players may choose from 3 sabers (one saber, dual sabers, staff), 3 saber styles (fast, medium, strong), 7 saber swings (left-to-right, right-to-left, bottom-left-to-top-right, bottom-right-to-top-left, top-to-bottom, top-right-to-bottom-left, top-left-to-bottom-right in each saber style), and 3 katas (one in each saber style). Players are able to move freely in 3D with the WASD keys and look freely with the mouse, even during a saber-swing. Players can also jump with space and switch saber styles with L. Swings are chosen with the direction keys. For instance a bottom-right-to-top-left swing corresponds to pressing the AS keys. Although published in 2003 Jedi Knight Jedi Academy still sees competitive play thanks to its deep combat system. An experienced player may be victorious 10-0 against a novice player. A professional player may be victorious 10-0 against an experienced player. We name our artificial intelligence Sapphire in reference to the dragon Sapphira from Eragon by Christopher Paolini, the Phi neural module from Pathak et al’s paper, and the gemstone skin of the blue shadowtrooper from the Jedi Knight game.

2 Reward-Based Learning

An actor A interacts with an environment E over discrete time steps. At each time step t the actor receives a view v_t of the environment and then makes an action a_t . The view v_t includes a reward r_t . Usually the actor’s goal is to maximize the sum of rewards R_t ,

$$R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots$$

$$\gamma \rightarrow^- 1$$

As our notation implies, we suggest starting off with a small γ and letting it grow, although we save this for future work.

3 Artificial Intelligence

Actors usually learn to improve their reward-maximization performance by, in one way or another, guessing the future sum of rewards R_t and then checking this guess at $t + 1$. The discrepancy between the guess and the check is what drives learning. In neural network-based deep learning, this discrepancy is minimized by nudging parameters in the neural networks which originally produced the guess. If the actor knows how to make perfect guesses as to which actions correspond to which future sum of rewards, then the actor knows how to play perfectly. An artificial intelligence is model-based if has a future prediction model.

4 Curiosity

An artificial intelligence is curious if part of its rewards correspond to errors in the future prediction model. Although there are many papers for 2D or toy 3D environments, there are relatively few papers for realistic (3D, realtime, competitive) games. Part of the reason may be because of the exploration challenge introduced by an additional dimension. In particular, this is the first time curiosity has been applied to a realtime 3D game that sees competitive play. Moreover, the artificial intelligence learns to defend itself against built-in bots in a few hours while maintaining a variety of play. In particular, while artificial intelligences without curiosity tend to stagnate and get stuck in a mode of behavior, a curious artificial intelligence maintains a variety of play.

5 Program

Our inspiration and implementation is based on Deepak Pathak et al’s Curiosity-driven Exploration by Self-supervised Prediction [1], Tim Pearce et al’s Counter-Strike Deathmatch with Large-Scale Behavioural Cloning [2], and Seungeun

Rho’s implementation [3] of A3C [4]. In particular, we use the curiosity module from [1], the interfacing and mouse movement bucketing ingenuity of [2] and the clean program of [3]. We are inspired by the couch-potato effect [5] to employ many embedding networks. We perform experiments without curiosity, with the curiosity module and with a modified smooth version of the curiosity module. Without curiosity the artificial intelligence easily gets stuck in some corner of the map and in a repetitive pattern of behavior. With curiosity the artificial intelligence maintains a variety of play and exploratory behavior, preventing it from ever getting stuck. With smooth curiosity the artificial intelligence also never gets stuck but also behaves in more human-like, less jittery way.

6 Motivation

Innate motivation is dedicated to unsupervised A3C learning. That is, instead of rewards coming from the environment the artificial intelligence is free to choose which goals it would like to undertake. The curiosity module not only rewards error in the future prediction model (which predicts the future state given the present state and action) but also employs an coordination model (which predicts behavior given a state transition). Moreover neither the future nor coordination models operate states directly but rather utilize a neural state compression module. The coordination model backpropagates errors to the neural state compression module. We modify the curiosity module of [1] to reward predictable actions thereby making AI behavior more human-like,

$$R_c = 100(E_f - E_a)$$

$$R_m = M$$

$$R_d = H + 0.5D$$

$$R_s = 100(V - 0.5L)$$

$$R_b = 1800$$

$$R = R_c + R_m + R_d + R_s + R_b$$

where R_c is the curiosity reward, E_f the future prediction error, and E_a the action prediction error, R_m the momentum reward, M the momentum (usually ranges from (0-1000)), R_d the health reward, H the player health, D the damage inflicted, R_s is the score reward, V is the number of victories, L is the number of losses, R_b is the baseline reward, and R is the combined reward which is received by the actor neural module. We use dropout instead of ϵ -greedy strategies.

7 Results

See videos. Sapphire runs on one 8GB graphics processing unit, explores half of Tattooine in half an hour, adequately defends itself against built-in bots, and

learns to do all this in a few hours. Human players say the AI is better and more interesting to duel against than built-in bots. The AI explores quickly even if $R = R_c$ and the input is only raw pixels. If R_c is set too small then the AI stagnates into one mode of behavior. Finally, we remark that there are awfully few papers in the literature which apply deep learning to realtime, competitive 3D games.

8 References

We thank the authors of these references for their curious dedication as well as friends and coworkers who helped with interfacing with the game and editing the paper,

- [1] Deepak Pathak, Pulkit Agrawal, Alexei A. Efros and Trevor Darrell. Curiosity-driven Exploration by Self-supervised Prediction. In ICML 2017.
- [2] Tim Pearce, Jun Zhu. Counter-Strike Deathmatch with Large-Scale Behavioural Cloning, In proc. of IEEE Conference on Games (CoG), Beijing, China, 2022.
- [3] Seungeun Rho. A3C. 2019.
- [4] Volodymyr Mnih, Adrià Puigdomènech Badia, Mehdi Mirza, Alex Graves, Timothy P. Lillicrap, Tim Harley, David Silver, Koray Kavukcuoglu. Asynchronous Methods for Deep Reinforcement Learning. arXiv, 2016.
- [5] Yuri Burda, Harrison Edwards, Amos Storkey, Oleg Klimov. Exploration by Random Network Distillation. arXiv, 2018.